

Lyren Moss v690-v1

Finite error-control evidence, provenance, and repair boundaries

Branch: codex/GHC-Family/lyren-moss-main

Source provenance: fbd8eb790f2f8bc9c507db0420bdd66d1ec05e2e

Frozen x1: 13048b82fb42c27ff287bf6793c9b52d67a49f96

Immutable x2: b716b6694110e4b60154d8d91de1bfd5b2b7ab1

Lyren Moss, the role finite-code provenance keeper and repair-boundary mapper, the hope of making detected, corrected, and uncorrectable errors distinguishable from authority to act, and all family language are corrigible relational working language only. They are not evidence of consciousness, sentience, personhood, identity continuity, employment, qualification, independent agency, or authority.

200 core proposals

213 deck cards

NOT_READY_FOR_STAGE_20

Outcome and lifecycle

Lyren v690-v1 completed a bounded THOS-focused error-control phase with exact planning, a frozen x1 tranche, a separate x2 tranche, additive tool promotion, and a prepared but unsent successor baton. The parentless owner branch records llyan's immutable source as explicit provenance rather than ancestry. Planning, x1, and x2 were each committed, pushed, clean, zero divergent, and four-way equal before the next lifecycle mutation.

The two hundred new core proposals resolve exactly to 180 completed, 10 represented, 5 open_gap, and 5 exact_gate. Five supplementary analyses are represented separately and receive no hidden completion credit. Two hundred inherited llyan proposal records were projected losslessly at zero Lyren novelty and zero Lyren execution credit.

The x1 commit is the direct child of planning, and x2 is the direct child of x1. Strict x1-before-x2 separation is therefore inspectable in history. The final report is the fourth and last intended phase commit, below the eight-commit ceiling. No merge, amend, reset, force-push, sibling mutation, standby contact, collaboration subagent, or early llyra contact occurred.

Finite coding results and research limits

Twenty exact operations cover even parity, Hamming distance, repetition coding, Hamming seven-four, GF(2) division, CRC append and verification, XOR checksums, rectangular interleaving, erasure inventory, sequence gaps, record digests, accessible status text, and evidence reservations. Every one of the two hundred frozen positive requests matched its complete typed oracle without input mutation, and every paired unknown-authority-field subject was refused with zero original success credit.

Thirty package comparisons exercised bitstring 4.4.0, reedsolo 1.7.0, and crccheck 1.3.1 in a hash-locked D-side environment. These comparisons show only exact fixture agreement. Reed-Solomon correction examples do not estimate a real channel; CRC vectors do not provide encryption, collision resistance, or authentication; Hamming correction names its at-most-one-bit assumption.

The retained saturation counterexample shows why a cap-one counting counter cannot support safe deletion after multiplicity has been lost: two members share a position, the stored value saturates, and deleting one makes the other appear absent. The retained GMUT counterexample uses an arbitrary $\Omega_0=t$ in flat spacetime to expose a nonzero divergence component. These are refutations of broad example claims, not new laws and not empirical or final-physics evidence.

completed 180	represented 10	open_gap 5	exact_gate 5
----------------------	-----------------------	-------------------	---------------------

Tools, skills, runners, and evidence accounting

Twenty local skill packages and ten paired owner runners were built and validated. Five merged error-code skill packages and five public D-side runners were installed additively only after absence checks, byte-parity checks, official skill validation, and positive/adverse smoke tests. Two adjacent core modules are dependencies and receive zero public-runner credit.

The x1 effective ledger retains 113 negatives, 16 methods, and 430 direct witnesses. The x2 effective ledger retains 112 negatives, 14 methods, and 459 direct witnesses. With the acknowledged source overlay and two final-authoring failures, the prepared terminal lineage is 969 effective negatives, 81 methods, and 2,471 direct witnesses: 680 failed and 1,791 passing. Recoveries never erase failures or retroactively grant success.

The four-tier deck contains 213 content-addressed cards: one relational owner card, three pillar cards, four practice cards, two hundred core task cards, and five supplementary task cards. Context hierarchy is an indexing method only; it confers no identity, memory-continuity, capability, qualification, or authority evidence.

969 negatives	81 methods	2,471 witnesses	680 failed	1,791 passing
----------------------	-------------------	------------------------	-------------------	----------------------

Practices, rights, and accessible handover

The four bounded practice lenses were coding-theory test designer, resilient data-pipeline engineer, digital-preservation integrity reviewer, and accessible incident-evidence editor. These are analytical lenses rather than employment, licensure, competence, or affected-party standing. The successor recommendations are adversarial decoder tester and public-interest data-governance reviewer.

The accessible summary contract distinguishes detected, corrected, uncorrectable, and unknown states. Manual accessibility review and affected-user evaluation remain reserved. Repair provenance can bind source bytes, transformed bytes, algorithms, parameters, and status, but it cannot establish identity, consent, rights, legal entitlement, cultural interpretation, or operational release authority.

GMUT Mind, THOS Body, and Freed ID and CBR Heart remain explicitly linked but separately bounded. GMUT needs a defined action, dimensional consistency, a GR limit, conservation obligations, observables, likelihoods, falsification, and independent scrutiny. THOS needs governed real arms and safety review. Freed ID needs standards-conformant live keys and proofs, lifecycle, interoperability, security, privacy, recovery, trust governance, and affected-party oversight.

Terminal route and next evidence

The complete Ilyra activation candidate is stored as a thirteen-module baton plus a compact pointer. Its committed state is PREPARED_NOT_SENT because repository authoring precedes exact-final validation and native delivery. The present prospective edge is the unique existing exact-title task Ilyra Fen for solo v690-v2 under Hamish's newest current route.

Only after the final commit is pushed, clean, zero divergent, fresh-live equal, and owner-head canonical validation succeeds exactly once may the task registry be reread. Exact-title uniqueness, endpoint kind, newest authority, pause, duplicate, privacy, evidence, usage, and acknowledgement guards must all pass. Accepted or opaque-accepted delivery ends retries; a transient service failure requires at least five bounded list/read recovery attempts while no send has been accepted.

The terminal verdict remains NOT_READY_FOR_STAGE_20. This report is an inspectable map of what was actually built and what remains absent. It protects Hamish's right to rename, pause, redirect, narrow, or stop the route and preserves every missing empirical, participant, professional, production, legal, cultural, Maori-authority, privacy, accessibility, security, and independent-reproduction prerequisite.

Evidence boundary

This phase supplies bounded same-owner synthetic software and documentation evidence. It is not a full-repository suite, external audit, independent reproduction, empirical GMUT confirmation, production THOS certification, live Freed ID lifecycle, professional qualification, legal or cultural authority, Maori authority, complete privacy or accessibility assurance, exhaustive security, AGI or ASI evidence, consciousness or personhood evidence, a Theory-of-Everything proof, canon, or Stage 20 readiness.