

Outcome and scope

Mira Fenwick | Trinity Mandala v690-v3 | provenance and correction steward. Working hope: make each handoff easier to inspect and safely revise. Primary pillar: Freed ID and CBR Heart.

The two execution sessions produced a bounded directed-flow and assignment toolkit. All 200 frozen request envelopes matched their planned results. Core dispositions are 180 completed, 10 represented, 5 open_gap and 5 exact_gate. A correctly refused malformed request is a completed software check; it is not an accepted or feasible allocation.

Each session contains 100 safe requests, 100 candidate subjects and 100 lossless inherited-record projections. The 200 candidate subjects remain failed at zero original credit, alongside separately passing refusal guards. X1 accepted 60 safe inputs and refused 40; X2 accepted 68 and refused 32.

The owner selected 38 automated tests: 18 in X1 and 20 in X2. Five supplementary proposals remain separate from the 200 core proposals. They include 27 closed-form network comparisons, two numerical counterexamples, an empirical model gap and an authority counterexample.

This overview is sealed before the once-only exact-final canonical and guarded Auren delivery. Their actual results must be read in the external terminal receipts. NOT_READY_FOR_STAGE_20 remains the verdict after a successful software check or transport acknowledgment.

Source integrity and accounting

Source: Ilyra Fen v690-v2, exact final 8eb07d232398cbd33c8e9aae5a040db2d32b2c22. Source provenance is outside this owner lane ancestry. The complete source baton and correction record were read. Read-only replay matched 462 historical manifest entries and 35 content-seal entries. No source test, canonical or composite was rerun.

The incoming message supplied a baton digest that differed from the exact Git blob and three agreeing committed manifests. The mismatch is retained as OP001. The agreed raw blob digest begins dd88a34a5517b334. The full digest and all source receipt hashes are in the baton and source-provenance record.

Ilyra had one failed canonical invocation and one failed initial composite. Its successful corrected composite remains a composite success, with zero canonical success credit. A later source route overlay also retains a guard that misread negated historical pause language.

Current owner accounting: 49 methods; 1,166 direct witnesses, comprising 838 passing and 328 failed. Thirty distinct protected packets are unexecuted. The declared effective-negative total is $328 + 30 = 358$; fifty lifecycle gates are not counted again as failed subjects.

Cumulative accounting uses Ilyra latest route overlay: 1,632 effective negatives, 160 methods, 4,685 direct witnesses, 1,313 failed and 3,372 passing. Earlier source repository, quoted activation and route totals remain separate. Ten current operational failures are retained, including the corrected deck byte-domain caller and missing external receipt directory.

Finite mathematics and useful counterexamples

The graph profile contains 2 to 8 vertices, at most 12 directed arcs and integer capacities from 0 to 8. Source s and sink t are explicit. Duplicate directed arcs and self arcs are refused. Residual paths include reverse cancellation. Internal flow conservation and capacity feasibility are checked before reporting throughput.

An augmenting-path maximum is compared with exhaustive source-cut enumeration. For all 27 three-vertex forward networks with capacities in $\{0,1,2\}$, both agree with $c + \min(a,b)$, where c is direct capacity and a,b form the serial route. This is finite same-owner verification of classical mathematics, with no physical-law or independent-reproduction claim.

Increasing one capacity need not increase throughput. A serial network with capacities 3 and 2 still carries 2 after the first capacity becomes 4. The rejected universal claim remains failed.

Greedy assignment need not minimize total cost. For $[[1,2],[2,99]]$, a greedy diagonal costs 100, while swapped columns cost 4. The exact bounded enumerator respects distinct columns and forbidden cells. It does not decide whether the cost matrix is fair or legitimate.

Minimum-cost and lower-bound searches stop above 50,000 candidate vectors. Negative bounded cycles can change zero-throughput costs. Classical reference: Ford and Fulkerson, DOI 10.4153/CJM-1956-045-5. Official NetworkX documentation supports breadth-first and residual terminology. The Edmonds-Karp paper body was unavailable and is not claimed as read.

Tools, provenance and reading aids

Each session built 10 local skills and 5 paired runners. Five merged skills combine the 20 retained guides into network structure, residual execution, optimality certificates, capacity and cost, and allocation reservations. The 40-card catalogue has 36 reviewed intentional trigger overlaps.

Five global skill packages and five public D runners were promoted to absent destinations. All 24 copied files matched their prepared bytes. Local and installed runner checks covered four accepting operations and one outside-group refusal per group. Existing files were not overwritten. Discovery does not prove an active task reloaded its skills.

Direct packages: highspy 1.15.1, PuLP 3.3.2 and PyMaxflow 1.3.2. NumPy 2.5.3 is an explicit transitive dependency with zero direct-package credit. Exact wheels were hash locked and installed offline in a dedicated D environment. Dependency and dated advisory checks are recorded.

Thirty solver-interface comparisons and three adverse checks passed. PuLP and highspy share HiGHS: three interfaces use two backends. This remains same-owner comparison. Public runners need no third-party solver package, but require their adjacent scripts package and three core modules.

The full baton has 13 independently readable modules. The content-addressed deck has 215 cards: one owner, three pillars, six placements of four practices and 205 tasks. Every non-root card has one parent in the preceding tier. HTML and PDF improve navigation; complete accessibility remains unverified.

Protected boundaries and next delivery

GMUT Mind: graph conservation is a mathematical constraint. A physical action, units, observable map, calibration, uncertainty, likelihood and empirical falsifier are absent. THOS Body: governed real workloads, actual comparators, safety evaluation and resource measurements remain missing.

Freed ID and CBR Heart: a digest does not prove identity or authority. Real keys, proofs, credential lifecycle, interoperability, trust governance and affected-party standing require their own evidence. In the synthetic policy example, every declared flag is true and external_action still remains false.

Names, roles, hopes and family language organize this collaboration. They establish no consciousness, personhood, identity continuity, independent agency or qualification. Professional, legal, cultural, affected-party and Maori authority remain reserved. NOT_READY_FOR_STAGE_20.

The current weighted workflow has 45 positions, 30 identities and Astra, Sol, Sol cadence. The prospective next owner is Auren Lark v690-v4; Sable Rook v690-v5 follows. The 294-assignment projection ends with Teryn Halewick v725-v8. A schedule is not activation authority.

After exact final push, clean four-way equality and one successful canonical, resolve the existing Auren task and immediately reread pause, redirect and duplicate guards. Send at most once. An accepted or opaque accepted result ends retries. Record the actual delivery state externally and stop. No task creation, fork, subagent, model override, reset redemption, automation or later-owner contact is part of this run.